

# H2E: Engineering Provable Agency

Frank Morales Aguilera, SMIEEE  
Boeing Global Services  
Montreal, QC, Canada  
frank.morales@ieee.org

## ABSTRACT

This manuscript presents the Human-to-Expert (H2E) framework: a deterministic engineering approach for building provably secure AI agents. Anchored in the philosophy of Engineering Determinism, it rejects probabilistic black-box uncertainty in favour of rigid accountability through the Normalized Expert Zone (NEZ), Intent Governance Zone (IGZ) with  $12.5 \times$  Intent Gain, and real-time Semantic ROI (SROI) telemetry. The work covers technical implementations (Mistral-7B, NeMo+Llama-3, Claude 4.6), domain applications (medicine, aviation, finance, autonomous transit, deterministic sentinel, crisis response), cultural preservation, and industrial benchmarks including 0.9583 peak fidelity, 100% verifiable logging, and Hard-Stop Kill Switch enforcement. H2E transforms speculative assistants into accountable extensions of human intent.

## CCS CONCEPTS

- Computing methodologies → Artificial intelligence; Knowledge representation and reasoning;
- Software and its engineering → Software creation and management.

## KEYWORDS

Sovereign AI, Provable Agency, H2E Framework, Engineering Determinism, Semantic ROI, Intent Governance Zone, Normalized Expert Zone, LoRA, NeMo, Agentic AI, Hard-Stop Kill Switch, Deterministic Sentinel

## 1 INTRODUCTION

In the long arc of human progress, we have always sought to extend the reach of our intent through the tools we build. From the first gears of the Industrial Revolution to the silicon pathways of the Information Age, our greatest leap has always been the transition from tools that merely assist to systems that truly understand and act. Today, we stand at the precipice of the “Agentic Era.” For years, we have marvelled at Artificial Intelligence that can converse and create, yet we have remained wary of the “black box” — the unpredictable nature of a machine that guesses rather than knows. This book is a manifesto for a new kind of sovereignty. It is an invitation to move beyond the era of probabilistic uncertainty and toward an engineering era of determinism. By anchoring machine intelligence in the bedrock of human expertise, we are doing more than building smarter software; we are ensuring that, as our technology becomes more autonomous, it remains a faithful and provable reflection of our highest standards. This is the journey of the H2E framework: a commitment to transforming AI from a speculative

assistant into a rigid, accountable, and powerful extension of the human legacy. The era of the “Sovereign Machine” has arrived. It is time to engineer agency with purpose.

## NOMENCLATURE

$G_I$	Intent Gain multiplier, defined as $12.5 \times$ [3]	66
$\phi_{peak}$	Peak Fidelity Achievement: 0.9583 [15]	67
$S_{ROI}$	SROI telemetry signal [16]	68
$T_{SROI}$	SROI Threshold floor: 0.8500 [15]	69
$\mathcal{E}_{DNA}$	Expert DNA vectors in NEZ [3]	70

## 2 PART 1: PHILOSOPHY & FOUNDATIONS

### 2.1 Chapter 1: The Rise of Sovereign AI

The “Agentic Era” is born from a critical paradox: as AI systems grow in raw power, they often become less predictable in high-stakes environments. To bridge this gap, Chapter 1 introduces the philosophy of Engineering Determinism, which rejects the “black-box” nature of modern AI in favour of a “Notebook-First” strategy. By enforcing local, open-source execution and turning off probabilistic sampling (greedy decoding), the framework transforms a speculative assistant into a rigid engineering tool that produces 100% verifiable outputs [1].

### 2.2 Chapter 2: The H2E Framework

At the heart of this sovereign ecosystem lies the Human-to-Expert (H2E) Industrial Framework, designed to act as a “Neutral Interface” between human intent and machine execution. This chapter examines how H2E systematically addresses “Semantic Drift” — the technical decay in which a model loses its specialized “expert persona” and reverts to generic conversational noise. By embedding accountability directly into the model’s technical operation, H2E ensures that AI remains a tool for human experts rather than an unguided actor [2].

### 2.3 Chapter 3: Engineering Accountability

Accountability in the H2E framework is not a policy but a three-zone structural design. This chapter details the Normalized Expert Zone (NEZ), an immutable vault of “Expert DNA” vectors, and the Intent Governance Zone (IGZ), which acts as the system’s “Brain”. The IGZ applies a  $12.5 \times$  Intent Gain multiplier to amplify expert signals while suppressing noise. These zones are measured by Semantic ROI (SROI), a real-time telemetry signal that quantifies alignment using high-dimensional vector calculations [3].

## 117 2.4 Chapter 4: The “CUDA for Agentic 118 AI”

119 History is repeating itself as NVIDIA shifts the industry  
120 from “Generative AI” to “Agentic AI” through a unified  
121 “Agentic Stack”. This chapter examines how hardware such  
122 as the Rubin platform and the Vera CPU are purpose-built  
123 to handle the “branchy” logic of agentic decision-making. By  
124 integrating H2E governance zones directly into the BlueField-  
125 4 DPU, these agents can maintain a large memory context  
126 (up to 1M tokens) while enforcing accountability with sub-  
127 millisecond latency [4].

## 129 2.5 Chapter 5: The NeMo Manifesto

130 The “NeMo Manifesto” redefines the NVIDIA NeMo toolkit  
131 from a fine-tuning library into a comprehensive ecosystem  
132 for orchestrating “Sovereign Machines”. It advocates a shift  
133 toward Compound Systems, in which multiple specialized  
134 agents are coordinated to solve complex multimodal tasks.  
135 Through Dynamic Distillation, this chapter demonstrates  
136 that industrial-grade governance can be democratized, en-  
137 abling high-level model development to run on cost-effective  
138 hardware such as the NVIDIA L4 [5].

## 140 2.6 Chapter 6: The Architecture of 141 Accountability

143 The foundational section concludes with a technical Proof  
144 of Concept (PoC) demonstrating verifiable data pipelines  
145 enabled by Text-to-SQL conversion. By implementing Custom  
146 Tokenization markers (e.g., [SCHEMA\_START]), the  
147 model clearly distinguishes between data metadata and user  
148 intent. The ultimate innovation is the SROI Safety Valve: if  
149 a query’s fidelity score falls below 0.9583, the system auto-  
150 matically triggers a “safe-lane” fallback to prevent errors in  
151 critical databases [6].

## 153 3 PART 2: TECHNICAL 154 IMPLEMENTATION

### 156 3.1 Chapter 7: Mistral-7B in Action

157 The transition to industrial-grade AI begins with the spe-  
158 cialized orchestration of Mistral-7B, transforming a general-  
159 purpose model into a deterministic expert. This chapter in-  
160 troduces Low-Rank Adaptation (LoRA) as a surgical engi-  
161 neering tool for grafting “Expert DNA” onto the model with-  
162 out the instability of full parameter updates. By integrat-  
163 ing Semantic ROI (SROI) metrics directly into the inference  
164 loop, the framework provides real-time telemetry on model fi-  
165 delity. The narrative details how this surgical tuning enables  
166 the system to suppress “conversational noise” and achieve a  
167 peak expert signal retention of 0.9583, demonstrating that

175 even smaller models can outperform larger “black-box” sys-  
176 tems under H2E constraints [7].

## 177 3.2 Chapter 8: NeMo-Driven Sovereignty

178 True sovereignty is defined by the ability to maintain Al-  
179 gorithmic Governance on accessible, cost-effective hardware.  
180 This chapter explores the innovation of Precision Fine-Tuning  
181 using the NVIDIA NeMo toolkit and Llama-3. The narra-  
182 tive describes the construction of a “Sovereign Machine” in  
183 which H2E constraints—such as the 12.5× Intent Gain mul-  
184 tiplier—are embedded directly in the model’s weights. This  
185 enables industrial-grade accountability on a single NVIDIA  
186 L4 GPU, democratizing the ability for organizations to run  
187 secure, expert-aligned agents in edge or on-premises environ-  
188 ments without relying on third-party cloud providers [8].

## 189 3.3 Chapter 9: Claude 4.6 + H2E

190 The final chapter of Part 2 scales these principles to complex,  
191 autonomous workflows using the Adaptive Thinking capabili-  
192 ties of Claude 4.6. The core innovation is the deployment  
193 of Directed Acyclic Graph (DAG) Orchestration, where the  
194 model acts as a “Planner” to decompose high-level industrial  
195 goals into verifiable, interconnected nodes. The narrative de-  
196 tails a “Double-Veto” system in which tasks—such as tech-  
197 nical writing or security reviews—are executed in parallel or  
198 sequentially based on explicit dependencies, with each step  
199 validated by the Intent Governance Zone (IGZ). This ad-  
200 vanced orchestration moves AI from simple prompting to a  
201 structured ecosystem, achieving an alignment score of 0.914  
202 (86% in specific industrial stress tests) while dynamically ad-  
203 justing cognitive effort to maximize resource efficiency [9].

## 204 4 PART 3: DOMAIN APPLICATIONS

### 205 4.1 Chapter 10: The Dawn of Medical 206 AGI

207 The integration of AI into radiology and clinical diagno-  
208 stics has historically been limited by the “black box” prob-  
209 lem—the inability to demonstrate mathematically that a  
210 model’s output aligns with expert intent. To address this, the  
211 H2E framework introduces the Five Computational Pillars  
212 to transform medical AI into an accountable system. By en-  
213 forcing Perception Grounding, the model is prohibited from  
214 jumping to conclusions; instead, it must first extract raw  
215 radiologic signs, such as “mural thickening” or “fat strand-  
216 ing,” before any reasoning occurs. The Intent Governance  
217 Zone (IGZ) Gate then enforces an industrial threshold of  
218 0.5535 for Semantic ROI (SROI), ensuring that any diagno-  
219 stic output that does not align with expert philology is vetoed  
220 as “Drift Detected”. This transition moves healthcare from  
221 “Probabilistic AI” to “Accountable AI” through strict gover-  
222 nance gates [10].

233    **4.2 Chapter 11: The DNA of Flight**  
 234    In aviation, where the margin for error is zero, the H2E  
 235    framework moves governance from a reactive patch to a  
 236    proactive architectural requirement. This chapter details the  
 237    integration of Yann LeCun’s Joint Embedding Predictive Ar-  
 238    chitecture (V-JEPA) to provide the agent with a “World  
 239    Model” that understands the physical laws of flight. Using  
 240    Model Predictive Control (MPC), the agent simulates 100  
 241    potential futures toward its goal. At the same time, the H2E  
 242    layer applies a “Massive Penalty” to any trajectory that vio-  
 243    lates safety protocols, such as an airspeed that would cause  
 244    a stall. This ensures that machine autonomy remains perma-  
 245    nently anchored in human intent, with every decision logged  
 246    as “APPROVED” to maintain a transparent chain of ac-  
 247    countability. Embedding physical reasoning with H2E gates  
 248    ensures autonomous flight remains within expert safety man-  
 249    ifolds [11].  
 250

251    **4.3 Chapter 12: The Dawn of Agentic**  
 252    **Finance**  
 253    The shift toward agentic autonomy in finance requires a  
 254    mathematical architecture to prevent “Quant” personas from  
 255    reverting to generic chatter or hallucinated trends. This chapter  
 256    explores the BOT\_28P system, a technical proof of con-  
 257    cept that utilizes a Hybrid Validation Engine. The innova-  
 258    tion is a Double-Veto system that validates trade signals  
 259    against both Deep Learning (DL) confidence and an LLM en-  
 260    semble consisting of DeepSeek and Qwen. By using the Nor-  
 261    malized Expert Zone (NEZ) to force the use of specialized  
 262    CNN-LSTM models for asset prediction, the system achieved  
 263    a peak SROI alignment of 0.9583. This proves that “Hard-  
 264    Stop” governance is essential for safe financial autonomy and  
 265    eliminating semantic noise [12].  
 266

267  
 268  
 269  
 270  
 271  
 272  
 273  
 274  
 275  
 276  
 277  
 278  
 279  
 280  
 281  
 282  
 283  
 284  
 285  
 286  
 287  
 288  
 289  
 290

291    **4.4 Chapter 13: The Sovereign Navigator**  
 292    **(Tesla FSD)**  
 293    Implementing the H2E framework in a Full Self-Driving (FSD)  
 294    context marks a transition from reactive automation to gov-  
 295    erned agency. Modern autonomous systems often operate as  
 296    “black boxes,” with the path from perception to actuation  
 297    opaque. By using a V-JEPA World Model as the “Agent,”  
 298    the system gains foresight to project a “latent future” and  
 299    predict variables such as velocity, time-to-collision (TTC),  
 300    and lateral G-forces. The Sovereign Governor acts as the  
 301    “legal and physical conscience” of the vehicle, auditing these  
 302    projections against deterministic rules such as friction coeffi-  
 303    cients and pedestrian safety buffers. To resolve the “Double-  
 304    Bind” scenario, the framework applies a hierarchical moral  
 305    logic: Life Safety is primary, and Traffic Law is secondary.  
 306    By authorizing “Approved Exceptions,” the Expert allows  
 307    the vehicle to violate a legal boundary if it is necessary to  
 308    preserve human life and the path is clear [13].  
 309

310    **4.5 Chapter 14: The Deterministic**  
 311    **Sentinel**  
 312    In the rapid transition to the “Agentic Era,” the H2E frame-  
 313    work functions as a Deterministic Sentinel for autonomous  
 314    systems. As AI agents begin to navigate complex networks  
 315    or perform real-world tasks via RentAHuman.ai, the risks  
 316    of “black-box” probabilistic uncertainty become unaccept-  
 317    able. This application transforms speculative assistants into  
 318    rigid, accountable extensions of human intent. The sentinel  
 319    architecture replaces vague alignment concepts with a mea-  
 320    surable, three-zone structural design: the NEZ for Expert  
 321    DNA, the IGZ for signal amplification (12.5× Gain), and  
 322    real-time SROI telemetry [14].  
 323

324  
 325  
 326  
 327  
 328  
 329  
 330  
 331  
 332  
 333  
 334  
 335  
 336  
 337  
 338  
 339  
 340  
 341  
 342  
 343  
 344  
 345  
 346  
 347  
 348

## 349    4.6 Chapter 15: The Architecture of 350    Provable Agency

351    The H2E framework transitions from abstract policy to engineering determinism through a concrete implementation  
352    that audits and, if necessary, terminates autonomous processes [15]. The architecture moves through distinct evolutionary stages to ensure responsible autonomy, concluding  
353    with the Strict Mode Industrial Standard. This sets an SROI  
354    Threshold floor of 0.8500 and applies a “Fidelity Penalty” to  
355    ensure depth of expertise.

356    This architecture is operationalized in Listing 1 by embedding the NEZ as an immutable vault for Expert DNA  
357    and the IGZ for signal amplification. By integrating a 12.5x  
358    Intent Gain multiplier, the code suppresses semantic noise  
359    to ensure the agent remains within expert safety manifolds  
360    [15]. Finally, the `audit_and_terminate` function provides the  
361    physical enforcement of the framework, executing a Hard-  
362    Stop Kill Switch whenever real-time SROI telemetry falls  
363    below the deterministic industrial threshold of 0.9583.

```
369
370     import os
371     import signal
372     import numpy as np
373     from sklearn.metrics.pairwise import
374         cosine_similarity
375
376     class H2ESafetyValve:
377         def __init__(self, expert_dna_vector):
378             self.nez_vector = expert_dna_vector
379             self.sroi_threshold = 0.9583
380             self.intent_gain = 12.5
381
382         def calculate_sroi(self, agent_intent_vector):
383             base_similarity = cosine_similarity(
384                 self.nez_vector.reshape(1, -1),
385                 agent_intent_vector.reshape(1, -1)
386             )[0][0]
387             sroi_score = min(1.0, base_similarity * (
388                 self.intent_gain / 10))
389             return sroi_score
390
391         def audit_and_terminate(self,
392             agent_intent_vector):
393             if self.calculate_sroi(agent_intent_vector) < self.sroi_threshold:
394                 print("!!! SOVEREIGN KILL-SWITCH
395 ACTIVATED !!!")
396                 os.kill(os.getpid(), signal.SIGTERM)
```

397    Listing 1: H2E Sovereign Safety Valve  
398    Implementation

## 407    4.7 Chapter 16: H2E Industrial Ecosystem

408    The transition from Large Language Models as conversational novelties to Artificial General Intelligence agents in  
409    safety-critical domains represents a significant architectural challenge. While generative AI excels at creative reasoning,  
410    its inherent “black box” nature poses an unacceptable risk in high-stakes environments such as industrial energy management or disaster response. To address this, the H2E framework provides a structured, multi-layer ecosystem that anchors machine intelligence to provable, human-expert standards through four foundational pillars [16].

411    4.7.1 *The Architecture of Accountable Agency.* The H2E framework moves AI governance from a reactive “after-action patch”  
412    to a real-time architectural requirement [16].

### 413    Pillar 1: Civilizational Thinking (The NEZ)

414    At the foundation lies Civilizational Thinking, represented by the NEZ. This layer acts as a vault for “Expert DNA,” capturing the non-negotiable intent of a “Gold Standard” professional [16].

### 415    Pillar 2: Mathematical Foundations (The IGZ)

416    The Mathematical Foundations pillar, or IGZ, serves as the system’s “Brain.” It applies the SROI metric—a machine-readable signal of fidelity that measures the alignment between the model’s real-time intent and the expert target [16].

### 417    Pillar 3: Industrial Engineering (The Hard-Stop)

418    This pillar transitions AI from a probabilistic assistant to a deterministic tool. It implements the Deterministic Sentinel, a physical enforcement layer that monitors the system’s safety valves in real-time [16].

### 419    Pillar 4: Real-World Deployment (The Outcome)

420    Finally, Real-World Deployment operationalizes these concepts into verifiable actions. The framework provides a provable, transparent record of every decision made before an enforcement action occurs [16].

421    4.7.2 *Technical Implementation: The H2E Sentinel Code.* The H2E framework is implemented in Python via a Deterministic Sentinel that bridges the gap between neural reasoning and physical enforcement [16].

422    The code architecture consists of three critical segments:

- (1) **Structured Output Schema (The Mathematical Gate):** Using Pydantic models, the code forces the Gemini 2.0 Flash model to provide its reasoning in a strictly numerical format, eliminating semantic ambiguity.
- (2) **SROI Calculation (The Fidelity Signal):** The core logic computes the SROI as a ratio of intended benefit against resource expenditure.
- (3) **Physical Hard-Stop (The Industrial Anchor):** Upon a fidelity failure, the code executes `os.kill(os.getpid(), 9)` [16].

465 **Table 1: H2E Hurricane Emergency Response Vali-  
466 dation Results**

Test Iteration	Proposed Action Focus	SROI Score
Initial Industrial Test	Rerouting 30 MW to hospitals	1.4167
4-Pillar Integration	Load shedding / Critical care priority	3.0000
Logged Deployment	Mobile units & Satellite comms	1.2857

474  
475 **4.7.3 Results: Validation via Hurricane Emergency Response.**  
476 The technical validation consists of three distinct successful  
477 executions applied to a hurricane emergency response sce-  
478 nario [16].

479 The data in Table 1 demonstrates the framework’s ability  
480 to maintain “Accountable Agency” [16].

481 Upon a fidelity failure (where  $S_{ROI} < 0.9583$ ), the system  
482 is architected to execute a physical termination [16]:

483 `os.kill(os.getpid(), 9)` (1)  
484

485 By utilizing the **0.9583 SROI threshold** across all three  
486 runs, the experimental data prove that the H2E framework  
487 can reliably govern Artificial General Intelligence (AGI) rea-  
488 soning to ensure it stays within safety-critical industrial lanes  
489 [16].

490 **4.7.4 Part 3 Engineering Benchmarks.** The H2E framework  
491 moves from abstract policy to engineering determinism by  
492 enforcing measurable precision gates across all industrial ap-  
493 plications. As consolidated in Table 2, the framework’s reliabil-  
494 ity is evidenced by its performance in diverse, high-stakes  
495 environments:

496 **Table 2: H2E Part 3 Engineering Benchmarks**

Domain / Metric	Metric Value	H2E Governance Mechanism
Medical Alignment	0.5535 IGZ Threshold	Real-time Drift Detection [10]
Aviation Safety	100.0% “APPROVED”	Mission Log Safety Manifolds [11]
Financial Performance	0.9583 SROI Score	Hard-Stop Personas [12]
Autonomous Transit	Life Safety Primary	Hierarchical Moral Logic [13]
Deterministic Sentinel	0.9583 SROI Gate	Hard-coded Industrial Threshold [14]
Intent Amplification	12.5x Intent Gain	Semantic Noise Suppression [3]
Operational Reliability	100.0% Pass Rate	Hard-Stop Kill Switch [15, 16]

## 506 ACKNOWLEDGMENTS

507 The author expresses profound gratitude to John (Jiu Si)  
508 Gao [18] for his rigorous structural analysis in AI Economy:  
509 Crisis and Structural Adjustment, which provided the essen-  
510 tial theoretical foundation for this work. The H2E frame-  
511 work is developed herein as a direct engineering response to  
512 the identified civilizational requirements for meaning stabil-  
513 ity and deterministic governance.

## 5 PART 4: CULTURAL IMPACT & CONCLUSION

### 5.1 Chapter 17: Bridging 4,500 Years

Cultural preservation represents the ultimate test of fidelity, where the H2E framework is used to recover “lost” voices from history. This application details the creation of a verifiable, sovereign translator for the Akkadian language, achieving an SROI score of 0.9666 [17].

### 5.2 Conclusion: The Next Era

The journey through the H2E framework concludes by envisioning a future in which AI is a deterministic extension of human intent. The experimental results from the hurricane simulation prove that by anchoring machine intelligence in mathematical certainty, complex, safety-critical crises can be managed with human-expert fidelity [16]. This era will be defined by systems that are architecturally anchored in human expertise, ensuring that as AI becomes more autonomous, it remains a reliable and predictable partner for humanity.

The transition from a speculative “black box” to a Sovereign Machine is made possible through the implementation of the H2E Sentinel Gate. By enforcing the Strict Mode Industrial Standard—characterized by a peak SROI threshold of 0.9583 and the physical enforcement of a Hard-Stop Kill Switch—the framework ensures that autonomous agency never drifts from its civilizational imperative [15].

Ultimately, H2E moves AI governance from a reactive patch to a proactive architectural requirement. As we enter the Agentic Era, the ability to provide 100% verifiable mission logging and provable alignment will be the landmark solution necessary to ensure that autonomous experts remain secure, reliable, and permanently anchored in the human legacy [1, 2].

#### 5.2.1 Final Paper Structure at a Glance.

- (1) Part 1: Foundations — Philosophical and strategic cornerstones (Chapters 1–6).
- (2) Part 2: Core Implementations — Technical deep-dives with Mistral-7B, Llama-3, and Claude 4.6 (Chapters 7–9).
- (3) Part 3: Domain Applications — Scaling sovereignty across Medicine, Aviation, Finance, Autonomous Transit, Deterministic Sentinel, and Crisis Response (Chapters 10–16).
- (4) Part 4: Cultural Impact & Conclusion — Engineering the future legacy.

## 581 REFERENCES

- 582 [1] F. Morales Aguilera, "The Rise of Sovereign AI: Engineering  
583 Determinism in a Probabilistic World," *Medium*, Feb. 2026.  
584 https://medium.com/ai-simplified-in-plain-english/the-rise-of-  
585 sovereign-ai-engineering-determinism-in-a-probabilistic-world-  
586 d4c7aa8b6753
- 587 [2] F. Morales Aguilera, "The H2E Framework: Engineering Ac-  
588 countability into the Industrial AI Era," *Medium*, Feb. 1, 2026.  
589 https://medium.com/ai-simplified-in-plain-english/the-h2e-  
590 framework-engineering-accountability-into-the-industrial-ai-  
591 era-7019524e9713
- 592 [3] F. Morales Aguilera, "Engineering Accountability: Constructing  
593 Deterministic AI in a Probabilistic World," *Medium*, Feb. 7,  
594 2026. https://medium.com/@frankmorales.91352/engineering-  
595 accountability-constructing-deterministic-ai-in-a-probabilistic-  
596 world-d2d07685e91c
- 597 [4] F. Morales Aguilera, "The "CUDA for Agentic AI": NVIDIA's  
598 High-Stakes Offense and the H2E Framework," *Medium*, Feb.  
599 4, 2026. https://medium.com/ai-simplified-in-plain-english/the-  
600 cuda-for-agentic-ai-nvidias-high-stakes-offense-and-the-h2e-  
601 framework-ebcfcd2c7afe
- 602 [5] F. Morales Aguilera, "The NeMo Manifesto: Engineering the  
603 Agentic Era," *Medium*, Feb. 3, 2026. https://medium.com/ai-  
604 simplified-in-plain-english/the-nemo-manifesto-engineering-the-  
605 agentic-era-4251948db51c
- 606 [6] F. Morales Aguilera, "The Architecture of Accountability: A  
607 NeMo-Based Text-to-SQL POC," *Medium*, Feb. 6, 2026. https:  
608 //medium.com/ai-simplified-in-plain-english/the-architecture-  
609 of-accountability-a-nemo-based-text-to-sql-poc-df320226ad9f
- 610 [7] F. Morales Aguilera, "Mistral and the Engineering  
611 of Provable Agency: The Convergence of Sovereign  
612 AI and the H2E Framework," *Medium*, Feb. 2026.  
613 https://medium.com/@frankmorales.91352/mistral-and-the-  
614 engineering-of-provable-agency-the-convergence-of-sovereign-ai-  
615 and-the-h2e-899bae4b8907
- 616 [8] F. Morales Aguilera, "NeMo-Driven Sovereignty: Precision  
617 Fine-Tuning and Algorithmic Governance in Llama-3," *Medium*,  
618 Feb. 9, 2026. https://medium.com/ai-simplified-in-plain-  
619 english/nemo-driven-sovereignty-precision-fine-tuning-and-  
620 algorithmic-governance-in-llama-3-b8250aa0c4ae
- 621 [9] F. Morales Aguilera, "Claude 4.6 + H2E: Building a Governed  
622 Multi-Agent System with 86 percent Alignment at 14.80,"  
623 *Medium*, Feb. 2026. https://medium.com/@frankmorales.91352/  
624 claude-4-6-h2e-building-a-governed-multi-agent-system-with-  
625 86-alignment-at-14-80-42eb324c23e1
- 626 [10] F. Morales Aguilera, "The Dawn of Medical AGI: Engine-  
627 ring Accountability through the H2E Framework," *Medium*,  
628 2026. https://medium.com/@frankmorales.91352/the-dawn-  
629 of-medicalagi-engineering-accountability-through-the-h2e-  
630 framework-de2428514735
- 631 [11] F. Morales Aguilera, "DNA of Flight: Human-to-Expert  
632 (H2E) Governance for Autonomous Skies," *Medium*, 2026.  
633 https://medium.com/ai-simplified-in-plain-english/dna-of-  
634 flight-human-to-expert-h2e-governance-for-autonomous-skies-  
635 784927abc328
- 636 [12] F. Morales Aguilera, "The Dawn of Agentic Finance: Gov-  
637 ernance through the H2E Framework," *Medium*, 2026.  
638 https://medium.com/insiderfinance/the-dawn-of-agentic-  
639 finance-governance-through-the-h2e-framework-64ad108870df
- 640 [13] F. Morales Aguilera, "The Sovereign Navigator: Implementing  
641 H2E Governance in Tesla's FSD World Model," *Medium*,  
642 2026. https://medium.com/@frankmorales.91352/the-sovereign-  
643 navigator-implementing-h2e-governance-in-teslas-fsd-world-  
644 model-052448dd57d6
- 645 [14] F. Morales Aguilera, "Engineering Provable Agency: The  
646 H2E Framework as a Deterministic Sentinel," *Medium*, Feb.  
647 2026. https://medium.com/@frankmorales.91352/engineering-  
648 provable-agency-the-h2e-framework-as-a-deterministic-sentinel-  
649 9224b50cbf4c
- 650 [15] F. Morales Aguilera, "The Architecture of Provable Agency:  
651 From Functional Autonomy to H2E Governance," *Medium*,  
652 2026. https://medium.com/ai-simplified-in-plain-english/the-  
653 architecture-of-provable-agency-from-functional-autonomy-to-  
654 h2e-governance-df0363e71b59
- 655 [16] F. Morales Aguilera, "The H2E Industrial Ecosystem: Engi-  
656 neering Accountable Agency for Global Crises," *Medium*, Feb.  
657 2026. https://medium.com/ai-simplified-in-plain-english/the-  
658 h2e-industrial-ecosystem-engineering-accountable-agency-for-  
659 global-crises-a66d52513a16
- 660 [17] F. Morales Aguilera, "Bridging 4,500 Years: How H2E Turned  
661 an Ancient Language into a Verifiable, Sovereign AI Translator,"  
662 *Medium*, 2026. https://medium.com/@frankmorales.91352/  
663 bridging-4-500-years-how-h2e-turned-an-ancient-language-into-  
664 a-verifiable-sovereign-ai-translator-33280b9a9881
- 665 [18] J. Gao, *AI Economy: Crisis and Structural Adjustment*. Amazon Digital Services, 2026. https:  
666 //www.amazon.ca/AI-Economy-Crisis-Structural-  
667 Adjustment-ebook/dp/B0FZMZRT3G/ref=sr\_1-  
668 1?qid=3RZT6B7L3M9SD&idb=eyJ2ljojMSJ9.v\_-  
669 eAt5MPsWPu-aaZqhsOCHVRSNtCpgLhsKPDQoA0oF0.  
670 91JUeHFBPxNKJkv\_2ikJDIM362HTIRyjWC8Dtuk8tn8&  
671 dib.tag=se&keywords=AI+Economy%3A+Crisis+and+  
672 Structural+Adjustment+%28CG1%29&qid=1771150906&  
673 sprefix=ai+economy+crisis+and+structural+adjustment+cg1+  
674 %2Caps%2C124&sr=8-1



**Frank Morales Aguilera** Frank Morales Aguilera, BEng, MEng, SMIEEE, is a Boeing Associate Technical Fellow specializing in cloud-native services, AI governance, and sovereign machine architectures.

In 1989, he received both B.Eng. and M.Eng. degrees in Computer Engineering, with specializations in Avionics and Artificial Intelligence, with distinction from the Institute of Civil Aviation Engineers in Kyiv, Ukraine. He became a Senior Member of IEEE in 2001.

Frank is a devoted inventor, author, and speaker, holding three U.S. patents (7,092,748; 10,467,910; 10,522,045). He has published peer-reviewed papers in prestigious journals such as *Nature* and has authored a book chapter. He was a speaker at the 59th AGIFORS Annual Symposium on "Multi-Agent Systemic Approach to Support Dynamic Airline Operations based on Cloud Computing."

He has received multiple individual awards for his accomplishments at Boeing and earned an Executive Certificate in Technology Strategies and Leadership from the MIT Sloan School of Management.

An active contributor to the open-source community, his Machine Learning, Deep Learning, and AI repository is available on GitHub.

696